

Parameterized Verification with Byzantine Model Checker

Igor Konnov

<igor@informal.systems>

Tutorial at FORTE, June 15, 2020

streaming from Vienna / Austria to Valletta / Malta

informal



Co-authors and contributors:



Helmut Veith



Ilina Stoilkovska



Marijana Lazić



Jure Kukovec



Francesco Spegni



Ulrich Schmid



Florian Zuleger



Nathalie Bertrand



Josef Widder



Annu Gmeiner



Roderick Bloem

Timeline



Fault-tolerant distributed algorithms and threshold automata

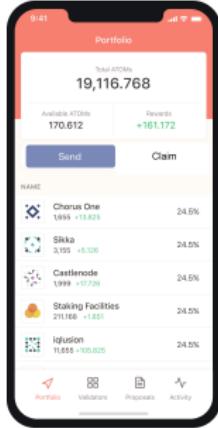
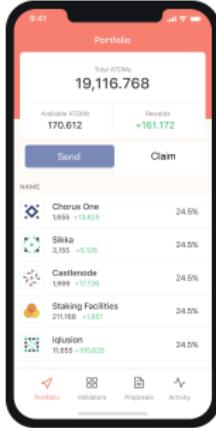


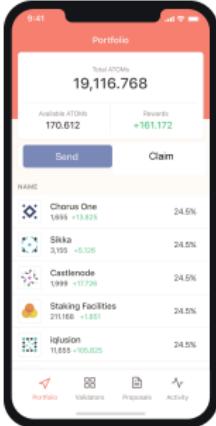
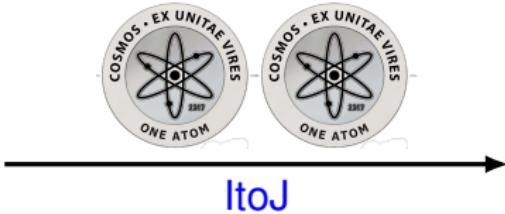
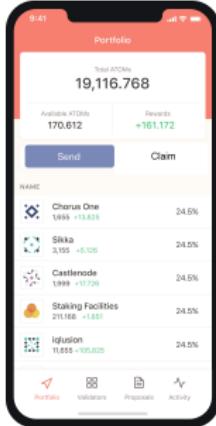
Safety of **asynchronous** threshold-guarded algorithms

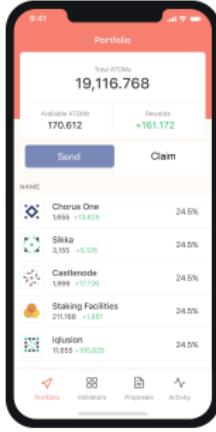


Liveness and **beyond** asynchronous algorithms

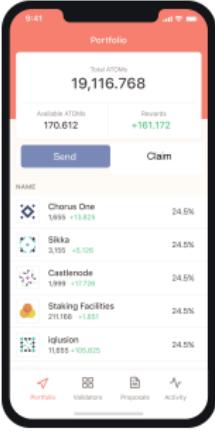
our inspiration:
distributed consensus





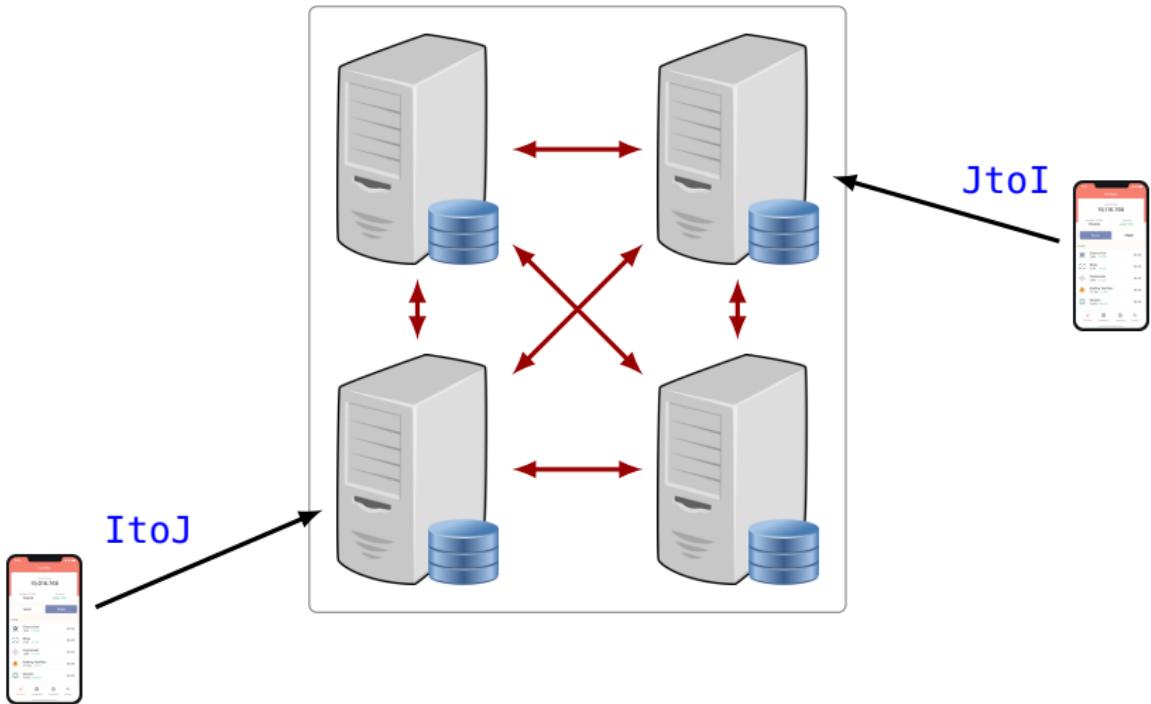


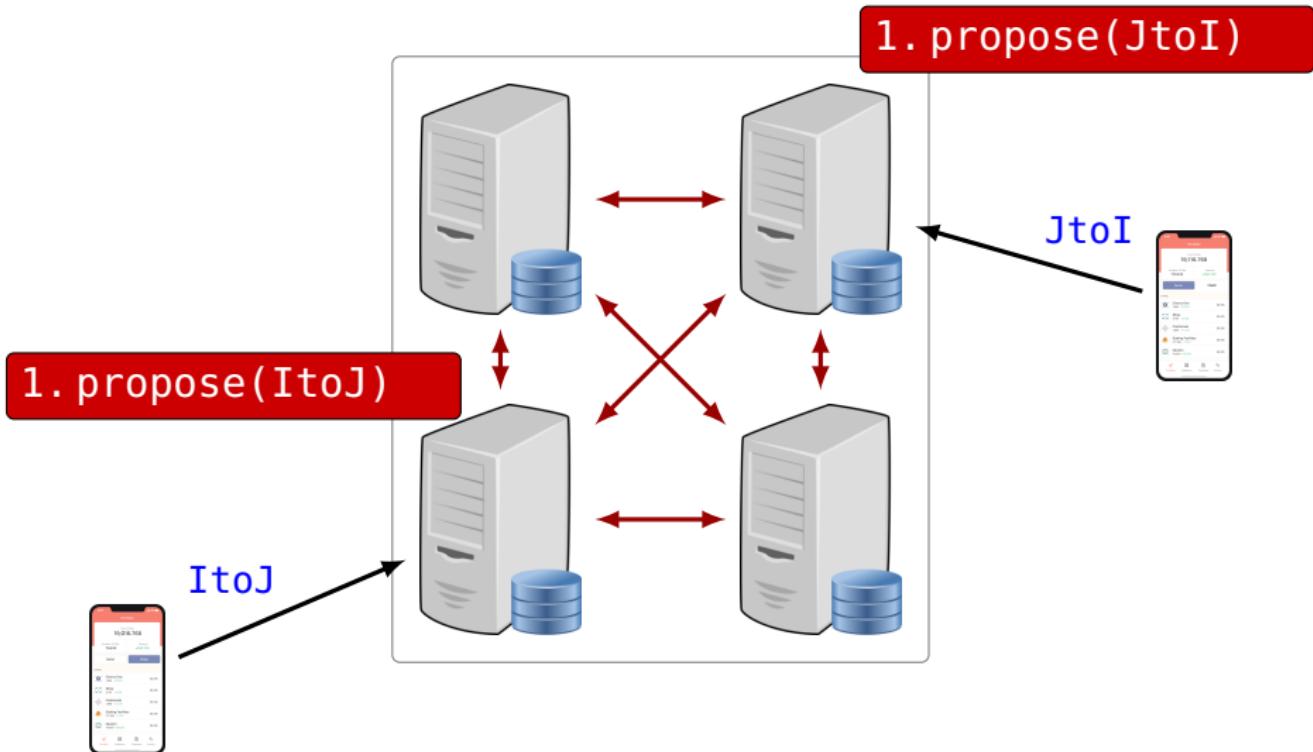
ItoJ

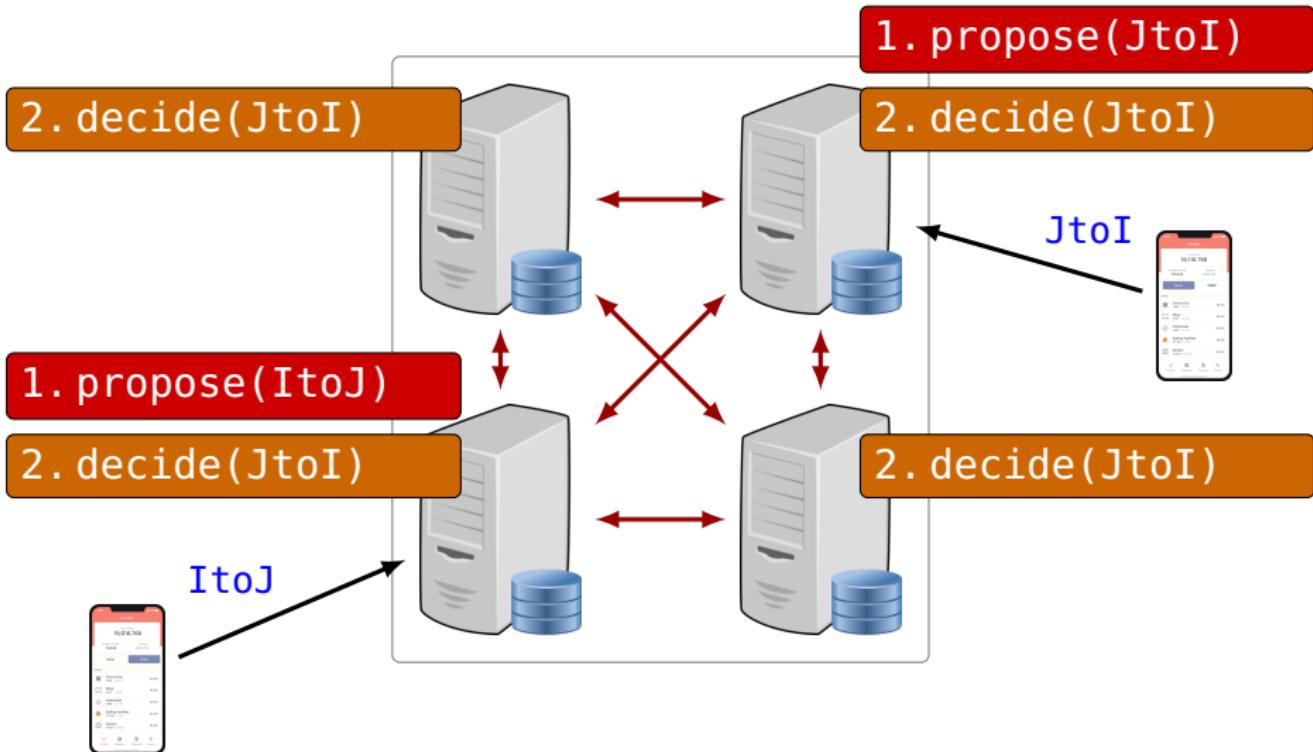


Jtol









Problem of Distributed Consensus

A distributed algorithm for n replicas
every replica proposes a value $w \in V$

Termination

every correct replica eventually decides on a value $v \in V$

Agreement

if a replica decides on v , no replica decides on $V \setminus \{v\}$

Validity

if a replica decides on v , the value v was proposed earlier

crucial to verify safety and liveness

Problem of Distributed Consensus

A distributed algorithm for n replicas
every replica proposes a value $w \in V$

Termination

every correct replica eventually decides on a value $v \in V$

Agreement

if a replica decides on v , no replica decides on $V \setminus \{v\}$

Validity

if a replica decides on v , the value v was proposed earlier

crucial to verify safety and liveness

~~Termination~~

every replica eventually decides on a value $v \in V$

Agreement

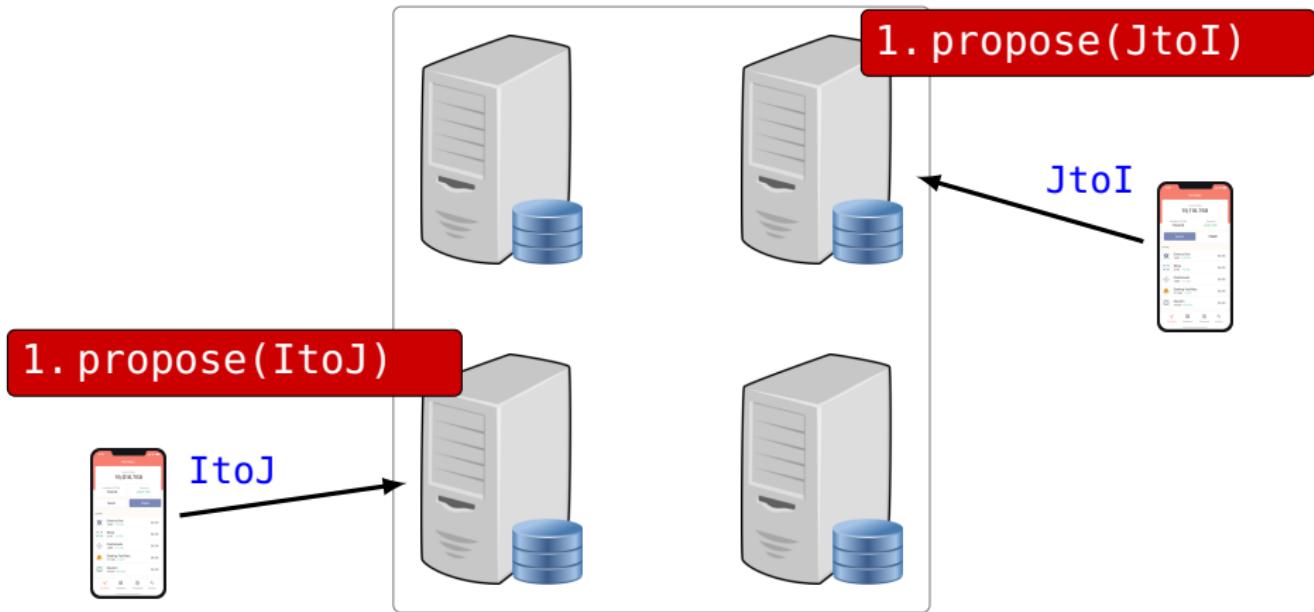
if a replica decides on v , no replica decides on $V \setminus \{v\}$

Validity

if a replica decides on v , the value v was proposed earlier

consensus without termination:

do nothing!



Naïve majority voting

n replicas follow the code:

- 1 input $u_i \in \{0, 1\}$
- 2 **send** u_i **to** all
- 3 **wait** until some value $v_i \in \{0, 1\}$ is **received** $\lceil \frac{n+1}{2} \rceil$ times
- 4 decide **on** v_j

Does it satisfy Validity, Agreement, and Termination?

What is the computation model?

n replicas follow the code:

- 1 input $u_i \in \{0, 1\}$
- 2 **send** u_i **to** all
- 3 **wait** until some value $v_i \in \{0, 1\}$ is **received** $\lceil \frac{n+1}{2} \rceil$ times
- 4 decide **on** v_j

Does it satisfy Validity, Agreement, and Termination?

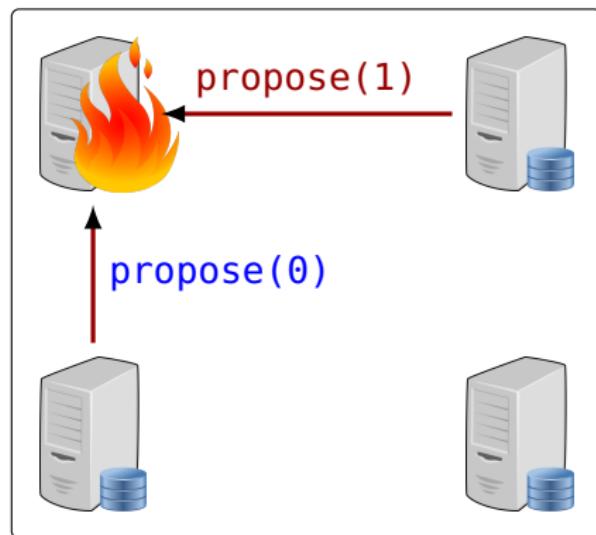
What is the computation model?

Asynchronous systems with faults

Various processor speeds

Various message delays, unbounded but finite

crashes



later today,

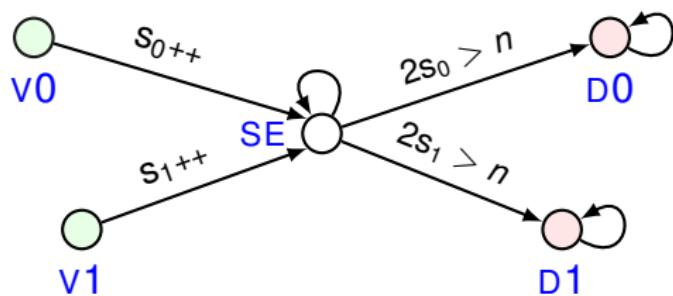
Byzantine

...

Formalizing pseudo-code...

- 1 input $u_i \in \{0, 1\}$
- 2 **send** u_i **to all**
- 3 **wait** until some value $v_i \in \{0, 1\}$ is **received** $\lceil \frac{n+1}{2} \rceil$ times
- 4 decide **on** v_j

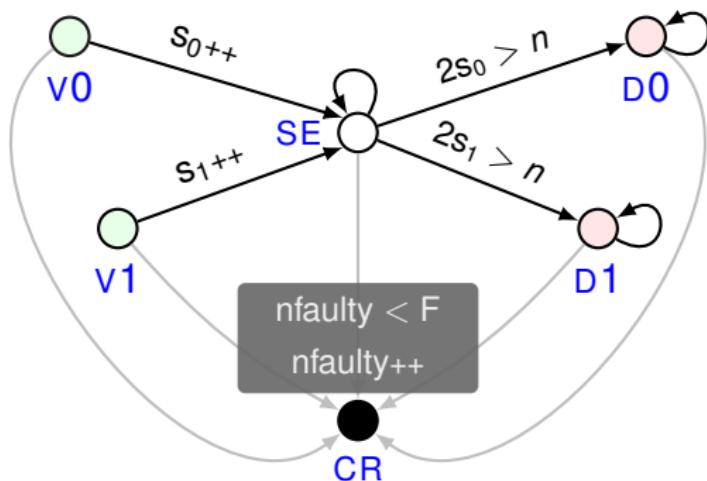
as a **threshold automaton**:



Formalizing pseudo-code...

- 1 input $u_i \in \{0, 1\}$
- 2 send u_i to all
- 3 wait until some value $v_i \in \{0, 1\}$ is received $\lceil \frac{n+1}{2} \rceil$ times
- 4 decide on v_j

as a threshold automaton:



Formalizing the distributed system...

replica 1: $u_1 = 0$ $v0 \xrightarrow{s_0++} SE$

replica 2: $u_2 = 1$ $v1 \xrightarrow{s_1++} SE$

replica 3: $u_3 = 0$

replica 4: $u_4 = 0$ $v0 \xrightarrow{s_0++} SE$

as a **counter system**:

$$\kappa_{v0} = 3 \quad \kappa_{v0} = 2$$

$$\kappa_{v0} = 1$$

$$\kappa_{v1} = 1 \quad \kappa_{v1} = 0$$

$$\kappa_{SE} = 0 \quad \kappa_{SE} = 1 \quad \kappa_{SE} = 2 \quad \kappa_{SE} = 3$$

$$s_0 = 0 \quad s_0 = 1$$

$$s_0 = 2$$

$$s_1 = 0 \quad s_1 = 1$$

...

Formalizing properties...

Termination: every replica eventually decides on a value $v \in V$

Agreement: if a replica decides on v , no replica decides on $V \setminus \{v\}$

Validity: if a replica decides on v , the value v was proposed earlier

as **temporal formulas**:

Termination: $\textit{fairness} \rightarrow \diamond (\kappa_{v0} = 0 \wedge \kappa_{v1} = 0 \wedge \kappa_{SE} = 0)$

Agreement: $\square (\kappa_{D0} = 0 \vee \kappa_{D1} = 0)$

0-Validity: $\kappa_{v1} = 0 \rightarrow \square (\kappa_{D1} = 0)$

1-Validity: $\kappa_{v0} = 0 \rightarrow \square (\kappa_{D0} = 0)$

Let's ask ByMC...

```
user@bymc: ~/fault-tolerant-benchmarks/forte20
user@bymc:~/fault-tolerant-benchmarks/forte20$ ./20x29
--limit-time: limit (in seconds) cpu time of subprocesses (ulimit -t)
--limit-mem: limit (in MB) virtual memory of subprocesses (ulimit -v)
-h|--help: show this help message

bymc options are as follows:
-O schema.tech=ltl           (default, safety + liveness as in POPL'17)
-O schema.tech=ltl-mpi        (parallel safety + liveness as in ISOLA'18)
-O schema.tech=cav15          (reachability as in CAV'15)
--smt 'lib2|z3|-smt2|-in'    (default, use z3 as the backend solver)
--smt 'lib2|mysolver|arg1|arg2|arg3' (use an SMT2 solver)
--smt 'yices'                (use yices 1.x as the backend solver, DEPRECATED)
-v                         (verbose output, all debug messages get printed)

Fine tuning of schema.tech=ltl:
-O schema.incremental=1 (enable the incremental solver, default: 0)
-O schema.noflowopt=1 (disable the control flow optimizations, default: 0
                      may lead to a combinatorial explosion of guards)
-O schema.noreachopt=1 (disable the reachability optimization, default: 0
                      i.e., reachability is not checked on-the-fly)
-O schema.noadaptive=1 (disable the adaptive reachability optimization, default: 0
                      i.e., the tool will not try to choose between
                      enabling/disabling the reachability optimization)
-O schema.noguardpreds=1 (do not introduce predicates for
                        the threshold guards, default: 0)
-O schema.compute-nschemas=1 (always compute the total number of
                           schemas, even if takes long, default: 0)
user@bymc:~/fault-tolerant-benchmarks/forte20$
```

Time for questions!



[bit.ly/2z8mE51]

(the examples and links for this talk)